

Scale Features and Build Model

Scales Raw Features

Import CSV of Aggregated Darshan Logs

Apply Log10 and Percent Scaling

```
In [1]: import os
import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
import random

from sklearn.model_selection import train_test_split
```

```
In [2]: df = pd.read_csv("./raws.csv",lineterminator='\n',sep = ',' ,error_bad_line
#df.mean())
```

```
In [3]: df = df.drop(df.columns[0],axis = 1)
df = df.drop(df.columns[0],axis = 1)
f = pd.DataFrame()
```

```
In [4]: df.columns
```

```
Out[4]: Index(['posix_read_time', 'posix_write_time', 'posix_meta_time',
'posix_bytes_read', 'posix_bytes_read_100', 'posix_bytes_read_1K',
'posix_bytes_read_10K', 'posix_bytes_read_100K', 'posix_bytes_read
_11M',
'posix_bytes_read_4M', 'posix_bytes_read_10M', 'posix_bytes_read_1
00M',
'posix_bytes_read_1G', 'posix_bytes_read_PLUS', 'posix_bytes_writ
e',
'posix_bytes_write_100', 'posix_bytes_write_1K',
'posix_bytes_write_10K', 'posix_bytes_write_100K',
'posix_bytes_write_1M', 'posix_bytes_write_4M', 'posix_bytes_write
_10M',
'posix_bytes_write_100M', 'posix_bytes_write_1G',
'posix_bytes_write_PLUS', 'posix_opens', 'posix_reads', 'posix_wri
tes',
'posix_seeks', 'posix_stats', 'posix_mmaps', 'posix_fsyncs',
'posix_fdsyncs', 'posix_rename_sources', 'posix_rename_targets',
'posix_renamed_from', 'posix_renamed_mode', 'posix_number_of_file
s',
'nprocs', 'posix_f_align', 'posix_m_align', 'lustre_number_of_file
s',
'lustre_mdts', 'lustre_osts', 'lustre_stripe_size',
'lustre_stripe_offset', 'lustre_stripe_width', 'lustre_number_of_o
sts',
'jobid', 'path'],
dtype='object')
```

In [5]:

```
df = df.dropna(axis=0, how='any')
df.columns
```

```
Out[5]: Index(['posix_read_time', 'posix_write_time', 'posix_meta_time',
              'posix_bytes_read', 'posix_bytes_read_100', 'posix_bytes_read_1K',
              'posix_bytes_read_10K', 'posix_bytes_read_100K', 'posix_bytes_read_
              _11M', 'posix_bytes_read_4M', 'posix_bytes_read_10M', 'posix_bytes_read_1
              00M', 'posix_bytes_read_1G', 'posix_bytes_read_PLUS', 'posix_bytes_writ
              e', 'posix_bytes_write_100', 'posix_bytes_write_1K',
              'posix_bytes_write_10K', 'posix_bytes_write_100K',
              'posix_bytes_write_1M', 'posix_bytes_write_4M', 'posix_bytes_write
              _10M', 'posix_bytes_write_100M', 'posix_bytes_write_1G',
              'posix_bytes_write_PLUS', 'posix_opens', 'posix_reads', 'posix_writ
              es', 'posix_seeks', 'posix_stats', 'posix_mmaps', 'posix_fsyncs',
              'posix_fdsyncs', 'posix_rename_sources', 'posix_rename_targets',
              'posix_renamed_from', 'posix_renamed_mode', 'posix_number_of_file
              s', 'nprocs', 'posix_f_align', 'posix_m_align', 'lustre_number_of_file
              s', 'lustre_mdts', 'lustre_osts', 'lustre_stripe_size',
              'lustre_stripe_offset', 'lustre_stripe_width', 'lustre_number_of_o
              sts', 'jobid', 'path'],
              dtype='object')
```

In [6]: #files

```
f['log10_p_files'] = df['posix_number_of_files'] + 1
f['log10_l_files'] = df['lustre_number_of_files']
```

In [7]: #accesses

```
df['p_accesses'] = df['posix_reads'] + df['posix_writes']
f['log10_p_accesses'] = df['p_accesses']

f['log10_p_accesses']
```

```
Out[7]: 0          0.0
1      880136.0
2     2379598.0
3     8903411.0
4     7846387.0
...
875282    2234152.0
875283     197651.0
875284         0.0
875285         0.0
875286    6065006.0
Name: log10_p_accesses, Length: 875287, dtype: float64
```

```
In [8]: #bytes
f['log10_p_bytes'] = df['posix_bytes_read']
```

```
In [9]: f['log10_p_opens'] = df['posix_opens']
f['log10_p_seeks'] = df['posix_seeks']
f['log10_p_stats'] = df['posix_stats']
f['log10_p_mode'] = df['posix_renamed_mode']
```

```
In [10]: f['log10_l_n_osts'] = df['lustre_number_of_osts']
f['log10_l_stripe_w'] = df['lustre_stripe_width']
f['log10_l_mdts'] = df['lustre_mdts']
```

```
In [11]: f['log10_p_nprocs'] = df['nprocs']
f['log10_p_falign'] = df['posix_f_align']
f['log10_p_malign'] = df['posix_m_align']
```

```
In [12]: f['perc_p_reads'] = df['posix_reads']
f['perc_p_writes'] = df['posix_writes']
```

```
In [13]: f['perc_p_bytes_read_100'] = df['posix_bytes_read_100']
f['perc_p_bytes_read_1K'] = df['posix_bytes_read_1K']
f['perc_p_bytes_read_10K'] = df['posix_bytes_read_10K']
f['perc_p_bytes_read_100K'] = df['posix_bytes_read_100K']
f['perc_p_bytes_read_1M'] = df['posix_bytes_read_1M']
f['perc_p_bytes_read_4M'] = df['posix_bytes_read_4M']
f['perc_p_bytes_read_10M'] = df['posix_bytes_read_10M']
f['perc_p_bytes_read_100M'] = df['posix_bytes_read_100M']
f['perc_p_bytes_read_1G'] = df['posix_bytes_read_1G']
f['perc_p_bytes_read_PLUS'] = df['posix_bytes_read_PLUS']
```

```
In [14]: f['perc_p_bytes_write_100'] = df['posix_bytes_write_100']
f['perc_p_bytes_write_1K'] = df['posix_bytes_write_1K']
f['perc_p_bytes_write_10K'] = df['posix_bytes_write_10K']
f['perc_p_bytes_write_100K'] = df['posix_bytes_write_100K']
f['perc_p_bytes_write_1M'] = df['posix_bytes_write_1M']
f['perc_p_bytes_write_4M'] = df['posix_bytes_write_4M']
f['perc_p_bytes_write_10M'] = df['posix_bytes_write_10M']
f['perc_p_bytes_write_100M'] = df['posix_bytes_write_100M']
f['perc_p_bytes_write_1G'] = df['posix_bytes_write_1G']
f['perc_p_bytes_write_PLUS'] = df['posix_bytes_write_PLUS']
```

```
f = f.replace(-np.inf, -1)
f = f.replace(np.nan, 0)
```

```
In [15]: df['time'] = df['posix_write_time'].astype('float') + df['posix_read_time']
```

```
In [16]: df['bytes'] = df['posix_bytes_read'].astype('float') + df['posix_bytes_writ
```

```
In [17]: #df = df[df['bytes'] > 99999999]
```

In [18]:

```
f['throughput'] = df['bytes'].astype('float') / df['time']

f = f[f['throughput'] > 0]
```

In [19]:

```
#delete columns with all zeros
f = f.loc[:, (f != 0).any(axis=0)]

#remove infinite values
f = f.replace([np.inf, -np.inf], np.nan).dropna(axis=0)

f.max()
```

Out[19]:

```
log10_p_files          1.219280e+05
log10_l_files          1.219270e+05
log10_p_accesses      2.251942e+10
log10_p_bytes         3.038456e+14
log10_p_opens         5.531094e+08
log10_p_seeks         1.445220e+10
log10_p_stats         6.522921e+07
log10_p_mode          5.337293e+07
log10_l_n_osts        3.600000e+02
log10_l_stripe_w      7.438575e+06
log10_l_mdts         1.000000e+00
log10_p_nprocs        3.520000e+05
log10_p_falign       1.422540e+11
log10_p_malign       1.085312e+06
perc_p_reads          2.237846e+10
perc_p_writes         1.302770e+10
perc_p_bytes_read_100 5.221517e+08
perc_p_bytes_read_1K  2.074657e+10
perc_p_bytes_read_10K 1.536278e+09
perc_p_bytes_read_100K 1.515506e+08
perc_p_bytes_read_1M  4.044503e+08
perc_p_bytes_read_4M  6.561462e+07
perc_p_bytes_read_10M 2.083200e+06
perc_p_bytes_read_100M 2.872090e+05
perc_p_bytes_read_1G  1.792000e+06
perc_p_bytes_write_100 1.302770e+10
perc_p_bytes_write_1K  2.852127e+09
perc_p_bytes_write_10K 3.867477e+08
perc_p_bytes_write_100K 8.347452e+07
perc_p_bytes_write_1M  1.357245e+07
perc_p_bytes_write_4M  3.839488e+06
perc_p_bytes_write_10M 6.190660e+05
perc_p_bytes_write_100M 1.249280e+06
perc_p_bytes_write_1G  1.937500e+04
throughput           2.344536e+09
dtype: float64
```

```
In [20]: t = pd.DataFrame()
t['throughput'] = f['throughput']
f = f.drop(labels = 'throughput', axis = 1)
f
```

Out[20]:

	log10_p_files	log10_l_files	log10_p_accesses	log10_p_bytes	log10_p_opens	log10_p_seeks
1	800.0	176.0	880136.0	2.390891e+10	8858.0	319241.0
2	361.0	224.0	2379598.0	5.019637e+10	62398.0	1107764.0
3	291.0	290.0	8903411.0	5.488943e+12	8711.0	2010273.0
4	320.0	201.0	7846387.0	2.293203e+10	23158.0	6015926.0
6	429.0	190.0	6647935.0	5.209185e+10	69261.0	4608438.0
...
875280	2.0	1.0	57344.0	3.006477e+10	1808.0	59128.0
875281	14.0	4.0	102439.0	7.467916e+08	193.0	91.0
875282	625.0	124.0	2234152.0	5.593977e+10	35112.0	1035457.0
875283	1089.0	1088.0	197651.0	1.465277e+09	2112.0	12509.0
875286	583.0	128.0	6065006.0	1.077412e+11	119432.0	2750807.0

671063 rows × 34 columns

```
In [21]: df = df[df.index.isin(t.index)]
t = t.reset_index()
f = f.reset_index()
f = f.drop(f.columns[0], axis = 1)
t = t.drop(t.columns[0], axis = 1)
```

```
In [22]: from sklearn.preprocessing import StandardScaler
f = StandardScaler().fit_transform(f)
```

In [23]: t

Out[23]:

	throughput
0	1.803194e+08
1	2.282342e+08
2	2.083669e+08
3	1.724841e+07
4	4.581690e+07
...	...
671058	5.482767e+08
671059	2.136637e+07
671060	1.362498e+08
671061	1.942413e+07
671062	1.842942e+08

671063 rows × 1 columns

In [24]: `print(t.min())`
`print(t.max())`

```
throughput    0.39201
dtype: float64
throughput    2.344536e+09
dtype: float64
```

In [25]: `rseed = 0`
`t_size = 0.1`

In [26]: `from sklearn.metrics import r2_score`
`from sklearn.metrics import mean_squared_error`
`from sklearn.metrics import mean_absolute_error`
`from sklearn.metrics import mean_squared_log_error`
`from sklearn.metrics import mean_absolute_percentage_error`

```
In [27]: from sklearn.linear_model import LinearRegression
for i in range(3):

    rseed = random.randint(1,10000)
    print(rseed)
    train_data, test_data, train_labels, test_labels = train_test_split(f,t

    reg = LinearRegression().fit(train_data, train_labels)

    predicted_labels = reg.predict(test_data)

    print("Mean True Value: \t",int(test_labels.mean() ))
    print("Mean Absolute Error: \t", int(mean_absolute_error(test_labels, p
    print("Mean Squared Error: ", mean_squared_error(test_labels, predicted
    print("Root Mean Squared Error: ", mean_squared_error(test_labels, pred
    print("MAPE : " + str(mean_absolute_percentage_error( test_labels, predi
    print("R2: " + str(r2_score(test_labels,predicted_labels)) + "\n")
```

466

```
Mean True Value:      84543501
Mean Absolute Error:   54153061
Mean Squared Error:   2.261908446988301e+16
Root Mean Squared Error: 150396424.39194825
MAPE :826.5616182370247
R2: -0.9317984212354609
```

55

```
Mean True Value:      84950034
Mean Absolute Error:   55005401
Mean Squared Error:   9147388791499650.0
Root Mean Squared Error: 95641982.369144
MAPE :1104.2638165971384
R2: 0.24535925719705243
```

8878

```
Mean True Value:      84309897
Mean Absolute Error:   54277374
Mean Squared Error:   9377959238697192.0
Root Mean Squared Error: 96839863.89239295
MAPE :1545.0408488673352
R2: 0.2049561655447404
```

```
In [28]: import xgboost as xg
```

In [29]: *#EXTREME GRADIENT BOOST*

```

for i in range(3):
    rseed = random.randint(1,10000)
    print(rseed)

    train_data, test_data, train_labels, test_labels = train_test_split(f,t
    xgb_r = xg.XGBRegressor(n_estimators = 1000, seed = 123)
    xgb_r.fit(train_data, train_labels)
    predicted_labels = xgb_r.predict(test_data)

    print("Mean True Value: \t",int(test_labels.mean() ))
    print("Mean Absolute Error: \t", int(mean_absolute_error(test_labels, p
    print("Mean Squared Error: ", mean_squared_error(test_labels, predicted
    print("Root Mean Squared Error: ", mean_squared_error(test_labels, pred
    print("MAPE : " + str(mean_absolute_percentage_error( test_labels, predi
    print("R2: " + str(r2_score(test_labels,predicted_labels)) + "\n")

```

```

1541
Mean True Value:      84338182
Mean Absolute Error:   16088329
Mean Squared Error:   1138614692905361.2
Root Mean Squared Error:  33743365.1686574
MAPE :114.38197797167567
R2: 0.9045144804003565

```

```

5641
Mean True Value:      84332413
Mean Absolute Error:   16384662
Mean Squared Error:   1181244012528314.0
Root Mean Squared Error:  34369230.607162476
MAPE :56.22594013095263
R2: 0.8992227778496157

```

```

8977
Mean True Value:      84289999
Mean Absolute Error:   16252869
Mean Squared Error:   1157630148155565.8
Root Mean Squared Error:  34023964.32157143
MAPE :51.37917112564205
R2: 0.9015652082234624

```

In []:


```

In [30]: from sklearn.tree import DecisionTreeRegressor

for i in range(3):
    rseed = random.randint(1,10000)
    print(rseed)

    train_data, test_data, train_labels, test_labels = train_test_split(f,t

    reg = DecisionTreeRegressor(max_depth = 4)

    reg.fit(train_data, train_labels)
    predicted_labels = reg.predict(test_data)

    print("Mean True Value: \t",int(test_labels.mean() ))
    print("Mean Absolute Error: \t", int(mean_absolute_error(test_labels, p
    print("Mean Squared Error: ", mean_squared_error(test_labels, predicted
    print("Root Mean Squared Error: ", mean_squared_error(test_labels, pred
    print("MAPE : " + str(mean_absolute_percentage_error( test_labels, predi
    print("R2: " + str(r2_score(test_labels,predicted_labels)) + "\n")

```

5626

```

Mean True Value:      84067576
Mean Absolute Error:   32864432
Mean Squared Error:   3195103903233775.0
Root Mean Squared Error: 56525250.1386219
MAPE :808.6107414878769
R2: 0.7303356522831486

```

3795

```

Mean True Value:      84232117
Mean Absolute Error:   32258290
Mean Squared Error:   3102852439300382.0
Root Mean Squared Error: 55703253.39960299
MAPE :1428.7644701294769
R2: 0.737042053264592

```

454

```

Mean True Value:      84362852
Mean Absolute Error:   33421235
Mean Squared Error:   3298164123928292.0
Root Mean Squared Error: 57429644.992184065
MAPE :3748.0546185631597
R2: 0.7236418495764236

```

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []:

In []: